

WEEKLY DATA ANALYSIS CHALLENGE #19

In [1]:

```
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
```

In [2]:

```
cst = pd.read_csv('coastal_db.csv')
cst.head()
```

Out[2]:

	coaster_name	Length	Speed	Location	Status	Opening date	Type	Manufacturer	Height restriction	Model	...	speed1	speed
0	Switchback Railway	600 ft (180 m)	6 mph (9.7 km/h)	Coney Island	Removed	June 16, 1884	Wood	LaMarcus Adna Thompson	NaN	Lift Packed	...	6 mph	9 km
1	Flip Flap Railway	NaN	NaN	Sea Lion Park	Removed	1895	Wood	Lina Beecher	NaN	NaN	...	NaN	Na
2	Switchback Railway (Euclid Beach Park)	NaN	NaN	Cleveland, Ohio, United States	Closed	NaN	Other	NaN	NaN	NaN	...	NaN	Na
3	Loop the Loop (Coney Island)	NaN	NaN	Other	Removed	1901	Steel	Edwin Prescott	NaN	NaN	...	NaN	Na
4	Loop the Loop (Young's Pier)	NaN	NaN	Other	Removed	1901	Steel	Edwin Prescott	NaN	NaN	...	NaN	Na

5 rows x 56 columns



Number of rows and columns

In [3]:

```
num_rows, num_columns = cst.shape
print(f'This dataset has {num_rows} rows and {num_columns} columns')
```

This dataset has 1087 rows and 56 columns

Is there any missing data?

In [4]:

```
cst.isnull().sum()
```

Out[4]:

coaster_name	0
Length	134
Speed	150
Location	0
Status	213
Opening date	250
Type	0
Manufacturer	59
Height restriction	256
Model	343

Height	122
Inversions	155
Lift/launch system	292
Cost	705
Trains	369
Park section	600
Duration	322
Capacity	512
G-force	725
Designer	509
Max vertical angle	730
Drop	593
Soft opening date	991
Fast Lane available	1018
Replaced	914
Track layout	752
Fastrack available	1068
Soft opening date.1	991
Closing date	851
Opened	1060
Replaced by	999
Website	1000
Flash Pass Available	1037
Must transfer from wheelchair	981
Theme	1043
Single rider line available	1006
Restraint Style	1065
Flash Pass available2	1041
Acceleration	1027
Restraints	1063
Name	1052
year_introduced	0
latitude	275
longitude	275
Type_Main	0
opening_date_clean	250
speed1	150
speed2	152
speed1_value	150
speed1_unit	150
speed_mph	150
height_value	122
height_unit	122
height_ft	0
Inversions_clean	0
Gforce_clean	725
dtype:	int64

In [5]:

```
remove = ['Length',
          'Speed',
          'Status',
          'Opening date',
          'Type',
          'Height restriction',
          'Model',
          'Height',
          'Inversions',
          'Lift/launch system',
          'Cost',
          'Trains',
          'Park section',
          'Duration',
          'Capacity',
          'G-force',
          'Designer',
          'Max vertical angle',
          'Drop',
          'Soft opening date',
          'Fast Lane available',
```

```

'Replaced',
'Track layout',
'Fastrack available',
'Soft opening date.1',
'Closing date',
'Opened',
'Replaced by',
'Website',
'Flash Pass Available',
'Must transfer from wheelchair',
'Theme',
'Single rider line available',
'Restraint Style',
'Flash Pass available2',
'Acceleration',
'Restraints',
'Name',
'latitude',
'longitude',
'speed1',
'speed2',
'speed1_value',
'speed1_unit',
'height_value',
'height_unit']

```

```

cst_2 = cst.drop(remove, axis=1)
cst_2.isnull().sum()

```

Out[5]:

```

coaster_name      0
Location          0
Manufacturer      59
year_introduced   0
Type_Main         0
opening_date_clean 250
speed_mph         150
height_ft         0
Inversions_clean  0
Gforce_clean      725
dtype: int64

```

In [6]:

```

cst_2 = cst_2.dropna()
cst_2.isnull().sum()

```

Out[6]:

```

coaster_name      0
Location          0
Manufacturer      0
year_introduced   0
Type_Main         0
opening_date_clean 0
speed_mph         0
height_ft         0
Inversions_clean  0
Gforce_clean      0
dtype: int64

```

In [7]:

```

num_row, num_column = cst_2.shape
print(f'This dataset has {num_row} rows and {num_column} columns')

```

This dataset has 282 rows and 10 columns

Display the summary statistics of the numeric columns using the describe method.

In [8]:

```
cst_2.describe()
```

Out[8]:

	year_introduced	speed_mph	Inversions_clean	Gforce_clean
count	282.000000	282.000000	282.000000	282.000000
mean	2000.109929	54.313121	1.758865	3.749433
std	15.370813	17.295847	2.342249	0.894507
min	1884.000000	6.000000	0.000000	0.800000
25%	1995.000000	45.000000	0.000000	3.400000
50%	2001.000000	54.050000	0.000000	3.950000
75%	2009.000000	65.000000	3.000000	4.300000
max	2022.000000	149.100000	14.000000	6.500000

Rename the following columns: { 'coaster_name' : 'Coaster_Name', 'year_introduce' : 'Year_Introduced', 'opening_date_clean' : 'Opening_Date', 'speed_mph' : 'Speed_mph', 'height_ft' : 'Height_ft', 'Inversions_clean' : 'Inversions', 'Geforce_celan' : 'Geforce' }

In [9]:

```
replace = {
    'coaster_name' : 'Coaster_Name',
    'year_introduced' : 'Year_Introduced',
    'opening_date_clean' : 'Opening_Date',
    'speed_mph' : 'Speed_mph',
    'height_ft' : 'Height_ft',
    'Inversions_clean' : 'Inversions',
    'Gforce_clean' : 'Geforce'
}
```

```
cst_2 = cst_2.rename(columns=replace)
cst_2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 282 entries, 0 to 1084
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Coaster_Name          282 non-null    object
1   Location               282 non-null    object
2   Manufacturer           282 non-null    object
3   Year_Introduced        282 non-null    int64
4   Type_Main              282 non-null    object
5   Opening_Date           282 non-null    object
6   Speed_mph              282 non-null    float64
7   Height_ft              282 non-null    object
8   Inversions             282 non-null    int64
9   Geforce                282 non-null    float64
dtypes: float64(2), int64(2), object(6)
memory usage: 24.2+ KB
```

Are there any duplicate rows?

In [10]:

```
duplicate = cst_2.duplicated().sum()
print(f'Total number of duplicate rows is {duplicate}')
```

Total number of duplicate rows is 0

What are the top 3 years with the most roller coaster

Tn [11]:

```
year_pivot = cst_2.pivot_table(index='Year_Introduced', values='Coaster_Name', aggfunc='count')
n = 3
top3_years = year_pivot.nlargest(n, 'Coaster_Name')
top3_years
```

Out[11]:

Coaster_Name	
Year_Introduced	
2000	14
2001	13
1999	12

What are the top 5 Locations with the most number of coasters

In [12]:

```
location = cst_2.pivot_table(index='Location', values='Coaster_Name', aggfunc='count')
n = 5
top10_location = location.nlargest(n, 'Coaster_Name')
top10_location
```

Out[12]:

Coaster_Name	
Location	
Other	12
Kings Island	10
Six Flags Magic Mountain	10
Alton Towers	9
Warner Bros. Movie World	6

What are the top 10 Manufacturers with the most number of coaster built

In [13]:

```
manufacturer = cst_2.pivot_table(index='Manufacturer', values='Coaster_Name', aggfunc='count')
n = 10
top10_manufacturer = manufacturer.nlargest(n, 'Coaster_Name')
top10_manufacturer
```

Out[13]:

Coaster_Name	
Manufacturer	
Bolliger & Mabillard	49
Vekoma	39
Intamin	28
Arrow Dynamics	23
Mack Rides	19
Gerstlauer	17
Rocky Mountain Construction	10

Anton Schwarzkopf Coaster_Name 8

Custom Coasters International 8

Premier Rides 8

What is the average speed? Also display a plot to show it's distribution.

In [14]:

```
avg_speed = cst_2['Speed_mph'].mean().round(2)
print(f'The average speed of roller coasters is {avg_speed} mph')
```

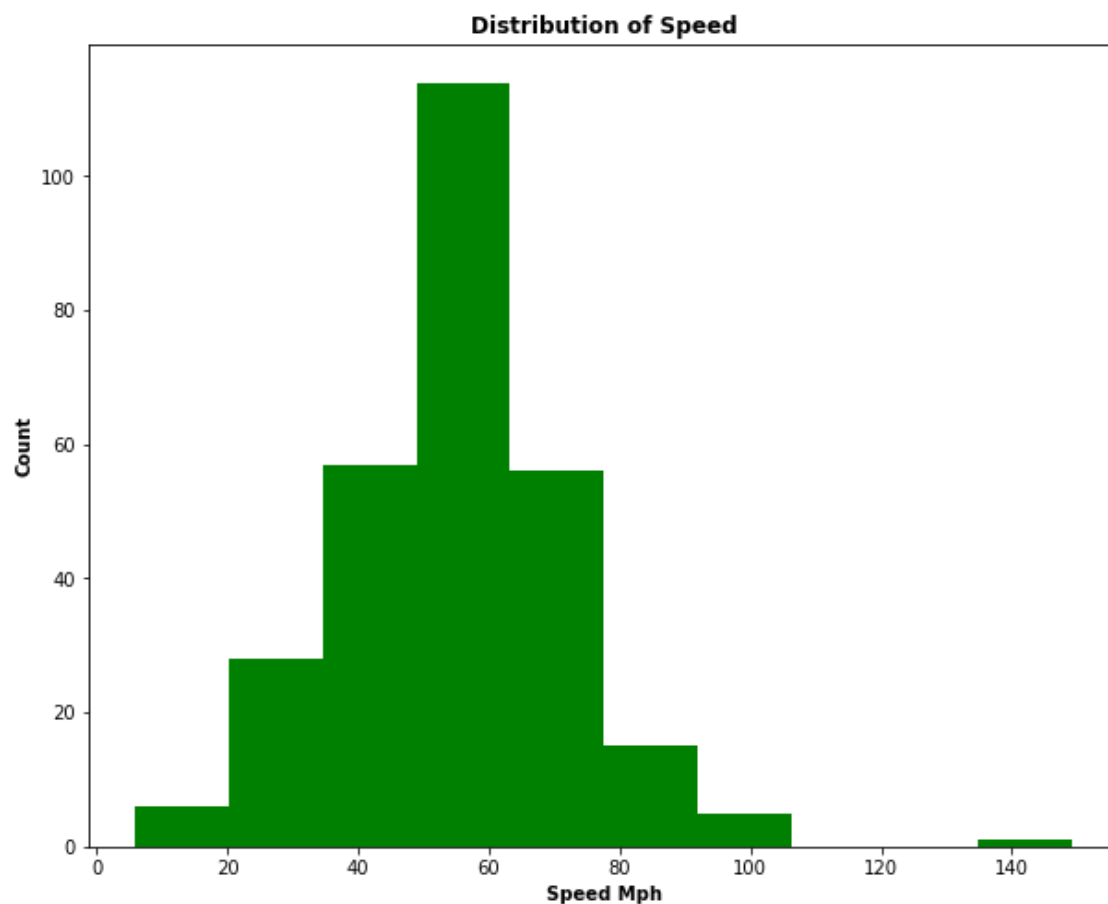
The average speed of roller coasters is 54.31 mph

In [15]:

```
fig = plt.figure(figsize=(10,8))
plt.hist(data=cst_2, x='Speed_mph', color='green')
plt.xlabel('Speed Mph', weight='bold')
plt.ylabel('Count', weight='bold')
plt.title('Distribution of Speed', weight='bold')
```

Out[15]:

Text(0.5, 1.0, 'Distribution of Speed')



Explore the feature relationships. Are there any positively or negatively correlated relationships

In [16]:

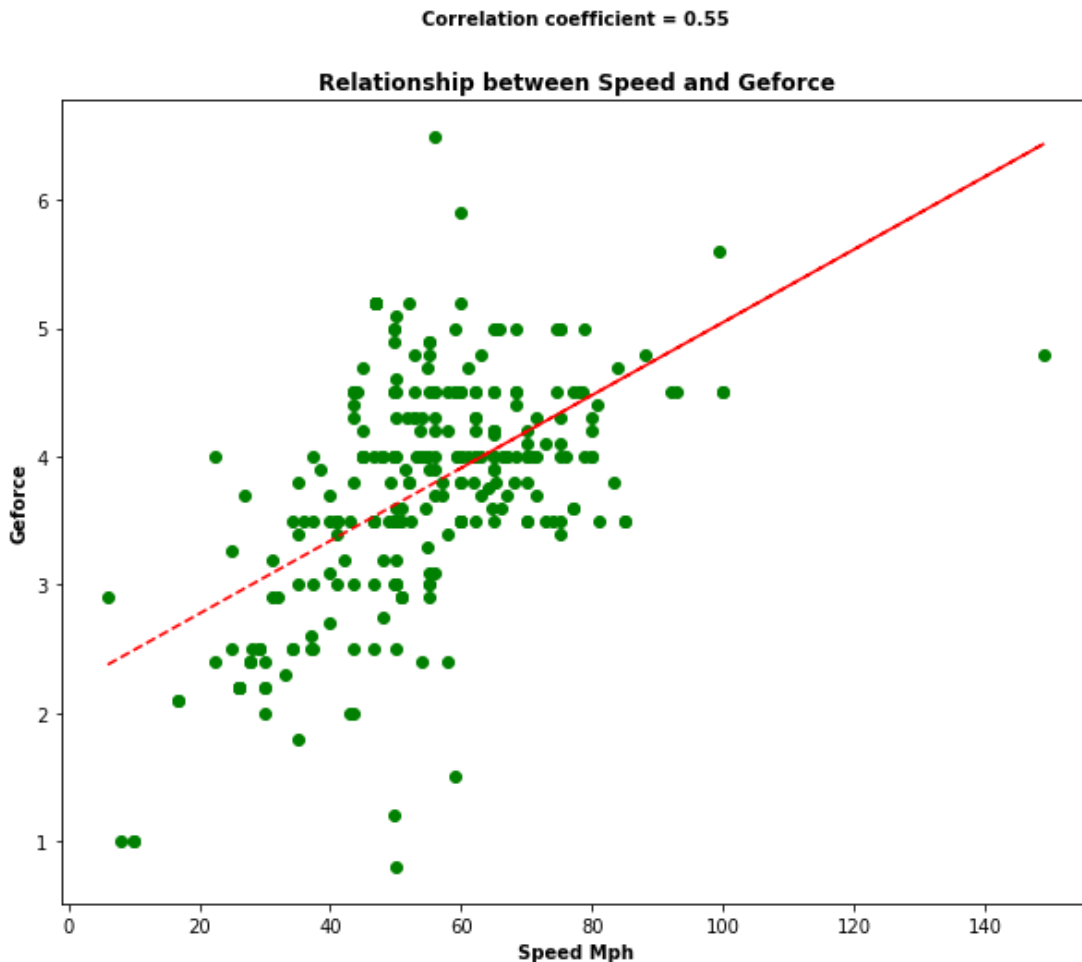
```
x = cst_2['Speed_mph']
y = cst_2['Geforce']

fig = plt.figure(figsize=(10,8))
plt.scatter(x, y, color='green')
z = np.polyfit(x, y, 1)
p = np.poly1d(z)
plt.plot(x, p(x), 'r--')
corr_coef = np.corrcoef(x, y)[1, 0]
```

```
plt.xlabel('Speed Mph', weight='bold')
plt.ylabel('Geforce', weight='bold')
plt.title('Relationship between Speed and Geforce', weight='bold')
plt.text(0.5, 1.1, f"Correlation coefficient = {corr_coef:.2f}", ha='center', va='center', weight='bold', transform=plt.gca().transAxes)
```

Out[16]:

Text(0.5, 1.1, 'Correlation coefficient = 0.55')



In [17]:

```
x = cst_2['Inversions']
y = cst_2['Geforce']

fig = plt.figure(figsize=(10,8))
sns.scatterplot(x, y, color='green')
z = np.polyfit(x, y, 1)
p = np.poly1d(z)
plt.plot(x, p(x), 'r--')
corr_coef = np.corrcoef(x, y)[1, 0]


plt.xlabel('Inversions', weight='bold')
plt.ylabel('Geforce', weight='bold')
plt.title('Relationship between Inversions and Geforce', weight='bold')
plt.text(0.5, 1.1, f"Correlation coefficient = {corr_coef:.2f}", ha='center', va='center', weight='bold', transform=plt.gca().transAxes)
```


c:\Python\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

Out[17]:

Text(0.5, 1.1, 'Correlation coefficient = 0.37')





In [18]:

```
try:
    cst_2['Height_ft'] = cst_2['Height_ft'].astype('float64')
except:
    cst_2['Height_ft'] = pd.to_numeric(cst_2['Height_ft'], errors='coerce')

cst_2['Height_ft'].info()
```

```
<class 'pandas.core.series.Series'>
Int64Index: 282 entries, 0 to 1084
Series name: Height_ft
Non-Null Count  Dtype
-----
207 non-null    float64
dtypes: float64(1)
memory usage: 4.4 KB
```

In [19]:

```
x = cst_2['Speed_mph']
y = cst_2['Height_ft']
z = cst_2['Geforce']

fig = plt.figure(figsize=(10,8))
sns.scatterplot(x, y, z, size=z, sizes=(20, 300))

plt.xlabel('Speed (mph)', weight='bold')
plt.ylabel('Height (ft)', weight='bold')
plt.title('Relationship between Speed, Height and Geforce', weight='bold')
```

c:\Python\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y, hue. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[19]:

```
Text(0.5, 1.0, 'Relationship between Speed, Height and Geforce')
```

Relationship between Speed, Height and Geforce

